

ISSUE1: SUSPECT MODE

Reason for suspect mode:

1. If one or more database files are not available.
2. If the entire database is not available.
3. If one or more database files are corrupted.
4. If a database resource is being held by the operating system.

How to recover?

➤ Scenario 1: If the file is full

Execute sp_resetstatus.

Syntax: sp_resetstatus database_name

Use ALTER DATABASE to add a data file or log file to the database.

Stop and restart SQL Server.

With the extra space provided by the new data file or log file, SQL Server should be able to complete recovery of the database

➤ Scenario2: If the data file was damaged.

- * Take T.Log backup
- * Restore last Full backup
- * Restore T.Log backup
- * Database comes online

➤ Scenario3: If the T.Log file was damaged

- * Take any user defined db for example: MyDB
- * Check the current location of files
sp_helpdb MyDB
- * Stop server
- * Move the T.Log file into different folder
- * Start server --> DB goes into suspect mode
Select databasepropertyex ('mydb','status')

Steps to Recover:

Step1: Make the db into single user

1) Alter database mydb set Single_User

Step2: Set the db into emergency mode

2) Alter database mydb set Emergency

Step3: Run checkdb with required repair level

3) DBCC CheckDB ('mydb', REPAIR_ALLOW_DATA_LOSS)

Step4: Set the db into multi user mode

4) Alter database mydb set Multi_User

ISSUE 2: MOVING MASTER DATABASE

1. Create two folders and grant read write permissions to service account

d:\master_data

e:\master_log

2. Find the current path

sp_helpdb master

3. Stop SQL Server

4. Move the files (master.mdf, mastlog.ldf) into new folders

5. Go to **SCCMSSCM** --> R.C on respective instance SQL Server Service -- properties --> Advanced --> Startup

Parameters--> Change the path of data and Log file

-dd:\master_data\master.mdf;-e....

-le:\master_log\mastlog.ldf

6. Apply --> OK

7. Start the service. Go to SSMS --> check the new path

sp_helpdb master

ISSUE 3: SHRINKDATABASE (DIFFERENT PROCESS)

If you ever want to transfer a large DB to a new one with more than one file, here is the way I am going to use (tested and approved)

1. Create a file which is as large as the data in your primary file (call it "buffer")
2. Empty the primary file (DBCC SHRINKFILE (<FILENAME>, EMPTYFILE))
3. Restart SQL Server Engine
4. Shrink the primary file to the Data size divided by the number of files you're gonna create (DBCC SHRINKFILE (<FILENAME>, NEWSIZE))
5. Create all the new files with the size of data divided by the number of files
6. Restrict their growth in order to fill the primary file in the next operation
7. Empty the buffer file (DBCC SHRINKFILE (BUFFER, EMPTYFILE))
8. Delete the buffer file (ALTER DATABASE REMOVE FILE (NAME=BUFFER))
9. Set final size of data files and ~~restrict~~~~unrestrict~~ their growth according to the final configuration needed

ISSUE 4: FIND OUT ~~FINDOUT~~ TABLE & INDEX SIZE

1. Create the temp table for further querying

CREATE TABLE #temp (

rec_id int IDENTITY (1, 1),

table_name varchar(128),

nbr_of_rows int,

data_space decimal(15,2),

index_space decimal(15,2),

total_size decimal(15,2),

percent_of_db decimal(15,12),

db_size decimal(15,2))

2. Get all tables, names, and sizes

EXEC sp_msforeachtable @command1="insert into #temp (no_of_rows, data_space, index_space) exec
sp_mstablespace '?'", @command2="update #temp set table_name = '?' where rec_id = (select max(rec_id) from
#temp)"

3. Set the total_size and total database size fields

```
UPDATE #temp SET total_size = (data_space + index_space), db_size = (SELECT SUM(data_space + index_space) FROM #temp)
```

4. Set the percent of the total database size

```
UPDATE #temp SET percent_of_db = (total_size/db_size) * 100
```

5. Get the data

```
SELECT *FROM #temp ORDER BY total_size DESC
```

6. Comment out the following line if you want to do further querying

```
DROP TABLE #temp
```

ISSUE 5: PAGE LEVEL RESTORE

```
SET NOCOUNT ON
```

```
USE master;
```

```
GO
```

```
CREATE DATABASE TestPageLevelRestore
```

```
ON
```

```
( NAME = TestPageLevelRestore,
```

```
  FILENAME = 'D:\TestPageLevelRestore.mdf',
```

```
  SIZE = 10)
```

```
LOG ON
```

```
( NAME = TestPageLevelRestore_log,
```

```
  FILENAME = 'D:\TestPageLevelRestore_log.ldf',
```

```
  SIZE = 5MB);
```

```
GO
```

```
Print 'Database TestPageLevelRestore Created'
```

```
ALTER DATABASE TestPageLevelRestore SET RECOVERY FULL
```

```
Print 'Recovery Model of database TestPageLevelRestore has been changed to FULL'
```

```
Use TestPageLevelRestore
```

```
GO
```

```
CREATE TABLE [Shift](
```

```
[ShiftID] tinyint IDENTITY(1,1) NOT NULL,
```

```
[Name] nvarchar(50) NOT NULL,
```

```
[StartTime] datetime NOT NULL,
```

```
[EndTime] datetime NOT NULL,
```

```
[ModifiedDate] datetime NOT NULL,
```

```
CONSTRAINT [PK_Shift_ShiftID] PRIMARY KEY CLUSTERED ([ShiftID] ASC)
```

```
)
```

```
Print 'Creation of Table "Shift" Completed'
```

```
SET IDENTITY_INSERT [Shift] ON
```

```
INSERT [Shift] ([ShiftID], [Name], [StartTime], [EndTime], [ModifiedDate]) VALUES (1, N'Day', '1900-01-01 07:00:00.000', '1900-01-01 15:00:00.000', '1998-06-01 00:00:00.000')
```

```
INSERT [Shift] ([ShiftID], [Name], [StartTime], [EndTime], [ModifiedDate]) VALUES (2, N'Evening', '1900-01-01 15:00:00.000', '1900-01-01 23:00:00.000', '1998-06-01 00:00:00.000')
```

```
INSERT [Shift] ([ShiftID], [Name], [StartTime], [EndTime], [ModifiedDate]) VALUES (3, N'Night', '1900-01-01 23:00:00.000', '1900-01-01 07:00:00.000', '1998-06-01 00:00:00.000')
```

```
SET IDENTITY_INSERT [Shift] OFF
```

```
Print 'Data Insertion to table "Shift" Completed'
```

```
BACKUP DATABASE TestPageLevelRestore TO DISK='D:\TestPageLevelRestore_FullBackup.bak' WITH STATS=10
Print 'Full Backup Completed'
```

```
--To get the list of index is id's from which you can choose one to corrupt
```

```
Use TestPageLevelRestore
```

```
Select * from sys.indexes where OBJECT_NAME(object_id)='Shift'
```

```
--To get the list of pages
```

```
DBCC IND ('TestPageLevelRestore', 'Shift',1)
```

```
-- To display the contents
```

```
DBCC TRACEON (3604);
```

```
GO
```

```
--TO get the page level data details
```

```
DBCC PAGE('TestPageLevelRestore',1,147,3);
```

```
--Get the Offset Value. This can be obtained by multiplying the page ID with 8192.
```

```
--Once you get the result copy the result and set the database to offline
```

```
SELECT 147*8192 AS [OffsetValue]
```

```
USE MASTER
```

```
ALTER DATABASE TestPageLevelRestore SET OFFLINE
```

Print 'Database TestPageLevelRestore is set ~~TestPageLevelRestore is set~~ to Offline. Now Open the TestPageLevelRestore.mdf file in the hex editor and press ctrl+g to go to thego the page where the index data is located.

Choose Decimal and paste the offset value.

once you go to the location, then manipulatemanuplate the value and save the file and exit the hexexit hex editor.

After manipulatingmanuplating data, bringdata-bring the databasedatabase online.'

```
--Run the below code after manipulating and exiting the hexexiting hex editor.
```

```
/*
```

```
USE MASTER
```

```
ALTER DATABASE TestPageLevelRestore SET ONLINE
```

```
Print 'Database TestPageLevelRestore is set to Online'
```

```
--Select the data and you will get error stating that the read failed at page (x:xxxx)
```

```
USE TestPageLevelRestore
```

```
Select * from shift
```

```
select * from sys.master_files where DB_NAME(database_id)='TestPageLevelRestore'
```

```
--Now Restore the page
```

```
USE master
```

```
-- Need to complete roll forward. So Backup the log tail.
```

```
BACKUP LOG TestPageLevelRestore TO DISK = 'D:\TestPageLevelRestore_log.bak' WITH INIT, NORECOVERY;
```

```
GO
```

```
Restore DATABASE TestPageLevelRestore Page='1:147' FROM DISK='D:\TestPageLevelRestore_FullBackup.bak'
```

```
-- restore the tail log backup.
```

```
RESTORE LOG TestPageLevelRestore FROM DISK = 'D:\TestPageLevelRestore_log.bak';
```

```
GO
```

```
--Verify
```

```
USE TestPageLevelRestore
```

```
Select * from shift
```

*/
SET NOCOUNT OFF

ISSUE 6: Identifying and Correcting the Transaction Log Full for User DB's

Error:

Subject: SQL Server Alert System: 'Full Log' occurred on CSQL2

Error: 9002, Severity: 17, State: 6

DESCRIPTION: The log file for database 'Dealix' is full. Back up the transaction log for the database to free up some log space.

Identifying:

- Configured (Log Full) alert to notify whenever there is a Transaction log Full on the User db.

And also Scheduled a job "SHRINKFILE" which performs the below tasks in the solution to prevent the Log Full. It is scheduled to run on Every Wednesday and ~~Sunday at 12 AM~~~~Sunday 12 AM~~ server time.

Solution:

If the Regular database Transaction logs runs out of space, this is indicated in the SQL ERRORLOG files, use the following process:

1. Free up (unallocated) the space used by the LOG portion of the database with the following command called from the master database:

USE master

GO

BACKUP LOG DBname WITH TRUNCATE_ONLY

GO

Note:

1. After you truncate a database LOG file, the SQL server documentation recommends that you ~~backup~~~~back-up~~ your database. In case of a physical failure (for example a power down or hard disk error), the SQL server cannot recover from the transaction log, as it was just truncated.

2. After running this command, the LDF file has been reorganized to have a lot of unallocated space, but the database must be *shrunk* to release that space to the file system. (It still looks like a large file if you view it from a command prompt directory listing). See next example for how to shrink the database.

Shrinking a Database

You can shrink a database to release the unallocated or unused space (or both) to the file system with the following command:

USE master

GO

DBCC SHRINKDATABASE (*database*)

GO

You can also use the SQL Enterprise Manager to shrink a database by selecting the following menu

Items: Right click on the Database -> All Tasks -> Shrink Database.

ISSUE 7: Troubleshooting host name changes

When the machine name is changed where we have installed SQL Server, all the instance~~instances~~ services are started but replication, Jobs, Alerts, Maintenance plans cause~~causes~~ errors. Hence we have to rename the instance.

To rename instance we can use the following SP

Steps:

1. Check the old server name as follows

SELECT @@servername

2. Drop the server and add the new server name

SP_DROPSERVER <oldName>

SP_ADDSERVER <newName>, local

3. Restart the instance

4. Check the server name again

SELECT @@servername

ISSUE 8: Troubleshooting User Connections

One of the users~~a user~~ is unable to connect to SQL Server. What may be scenarios and how to troubleshoot it?

Possible Scenarios

1. Error: 26

* SQL Browser

* Firewall

* No connectivity between client and server

2. Error: 28

* Instance TCP/IP was disabled

3. Error: 40

- * Instance service is not running

4. Error: 18456

- * Login failed. (Invalid login or pwd)

5. Expired Timeout

- * Network issue
- * Server is busy
- * In server max sessions are open
- * No available session memory

6. Connection Forcibly Closed

Update the client computer to the server version of the SQL Server Native Client.

7. In single user mode if any other service is connected with the db Engine, it doesn't allow connections.

ISSUE 9: Troubleshooting SQL Server Service Problems

My SQL Server service ~~has not~~ started. What may be the possible scenarios?

Possible Scenarios

- * Logon Failure
- * Problem with service account.
- * 3417
 - * Files are not present in the respective path or there are no permissions ~~on the target~~ folder where the files are not present.
- * 17113
 - * Master files are moved to different ~~locations~~, but not mentioned in startup parameters.
- * Service cannot be started in timely fashion
- * Insufficient resources, try to stop some other instances and start again.

How to find ~~error~~?

1. Using windows event log

- * start --> run --> eventvwr
- * System
- * In the right side check for the errors

2. Using SQL Server ErrorLog file

- * Go to respective instance LOG folder and open ErrorLog in notepad and check for the errors.

ISSUE 10: Troubleshooting database suspect mode-17207

Database has gone into suspect mode. How to handle this scenario?

Possible Scenarios

- * If the database files are corrupted or there ~~is a disk~~ issue.

* If restoration fails.

* ~~If the data~~~~file~~ ~~data~~ file was full.

Steps:

1. Check the error log for possible reasons.
2. ~~If a data~~~~file~~ ~~data~~ file was damaged or disk failure ~~happened~~~~happen~~.
 - Take tail log backup.
 - Restore full backup
 - Restore recent differential backups
 - Restore all log backups if any, made after recent differential backup.
 - Restore tail log backup WITH RECOVERY.
3. ~~If a log~~~~file~~ ~~log~~ file was damaged or disk failure ~~happened~~~~happen~~.
 - Try to take tail log backup with another copy ~~of the log~~~~of log~~ file if available with RAID level.
 - If the log file is not available then make it online by running the following commands where there may be data loss.

--step1: Make the db into single user

1) ALTER DATABASE <databaseName> SET Single_User

--step2: Set the db into emergency mode

2) ALTER DATABASE <databaseName> SET Emergency

--step3: Run checkdb with required repair level

3) DBCC CHECKDB ('<databaseName>', REPAIR_ALLOW_DATA_LOSS)

--step4: Set the db into multi user mode

4) ALTER DATABASE <databaseName> SET Multi_User

ISSUE 11: Troubleshooting master database corrupted

One of ~~the instance~~~~an instance~~ master database data ~~files~~~~file~~ was corrupted and I was unable to start the server. How to troubleshoot this scenario?

Possible Scenarios

If the master files are corrupted or damaged, ~~an instance~~~~instance~~ cannot be started. We have to ~~rebuild the master~~~~rebuild-master~~ database then by running the server in single user mode we have to ~~restore the latest~~~~restore latest~~ backup to get previous settings.

Steps

1. Check the error log ~~for the exact~~~~for exact~~ reason.
2. Rebuild master database as follows by running setup from
C:\Program Files\Microsoft SQL Server\100\Setup Bootstrap\Release
For windows authentication:

```
setup /ACTION=REBUILDDATABASE /QUIET /INSTANCENAME=<instance  
name> /SQLSYSADMINACCOUNTS=<accounts>
```

For mixed mode:

setup /ACTION=REBUILDDATABASE /QUIET /INSTANCENAME=<instance name> /SQLSYSADMINACCOUNTS=<accounts> /SAPWD=password

3. Once rebuilding is completed then run the server in single user mode
4. ~~Restore the master~~Restore master database by ~~replacing the existing~~replacing existing one.
5. Restart the server in multi user mode.

ISSUE 12: Troubleshooting Tempdb moving issue

One of ~~the instance~~an instance master database data ~~files~~file was corrupted and I was unable to start the server. How to troubleshoot this scenario?

Possible Scenarios:

* If the master files are corrupted or damaged, ~~an instance~~instance cannot be started. We have to ~~rebuild the master~~rebuild master database then by running the server in single user mode we have to ~~restore the latest~~restore latest backup to get previous settings.

Steps:

1. Check the error log ~~for the exact~~for exact reason.
2. Rebuild master database as follows by running setup from
C:\Program Files\Microsoft SQL Server\100\Setup Bootstrap\Release

For windows authentication:

setup /ACTION=REBUILDDATABASE /QUIET /INSTANCENAME=<instance name> /SQLSYSADMINACCOUNTS=<accounts>

For mixed mode:

setup /ACTION=REBUILDDATABASE /QUIET /INSTANCENAME=<instance name> /SQLSYSADMINACCOUNTS=<accounts> /SAPWD=password

3. Once rebuilding is completed then run the server in single user mode
4. ~~Restore the master~~Restore master database by ~~replacing the existing~~replacing existing one.
5. Restart the server in multi user mode.

ISSUE 13: Troubleshooting Insufficient Disk Space in tempdb

Running out of disk space in tempdb can cause significant disruptions in the SQL Server production environment and can prohibit applications that are running from completing operations.

Possible Scenarios:

Error	Is raised when
1101 or 1105	Any session must allocate space in tempdb.
3959	The version store is full. This error usually appears after a 1105 or 1101 error in the log.
3967	The version store is forced to shrink because tempdb is full.

3958 or 3966	A transaction cannot find the required version record in tempdb.
--------------	--

ISSUE 14: ~~FIND OUT~~ FIND OUT DATABASE SIZE

```
SELECT alt.filename [File Name] ,alt.name [Database Name] ,alt.size * 8.0 / 1024.0 AS [Originalsize (MB)] ,files.size * 8.0 / 1024.0 AS [Currentsize (MB)] FROM master.dbo.sysaltfiles alt INNER JOIN dbo.sysfiles files ON alt.fileid = files.fileid WHERE alt.size <> files.size
```

The above query allows us to find the current status of our databases and their corresponding final file growths. Use further filter conditions to fetch the databases that are of interest to you.

ISSUE 15: IMPORTANT SCRIPTS IN ~~SQL DBA~~SQLDBA

1. To display version, level, edition etc

```
select SERVERPROPERTY('productversion'),
        SERVERPROPERTY('productlevel'),
        SERVERPROPERTY('edition'),
        SERVERPROPERTY('isclustered')
```

2. To display execution plans present in procedure cache

```
SELECT cp.objtype AS PlanType,
        OBJECT_NAME(st.objectid,st.dbid) AS ObjectName,
        cp.ref countcp.refcounts AS ReferenceCounts,cp.use countcp.usecounts AS UseCounts,
        st.text AS SQLBatch,qp.query_plan AS QueryPlan
FROM sys.dm_exec_cached_plans AS cp
CROSS APPLY sys.dm_exec_query_plan(cp.plan_handle) AS qp
CROSS APPLY sys.dm_exec_sql_text(cp.plan_handle) AS st;
```

3. To display instance names with T-SQL

EXECUTE xp_regread

@rootkey ='HKEY_LOCAL_MACHINE',

@key ='SOFTWARE\Microsoft\Microsoft SQL Server',

@value_name ='InstalledInstances'

4. Backups Information

Note: To check recent backups of all databases

SELECT T1.Name AS DatabaseName,

COALESCE(CONVERT(VARCHAR(12), MAX(T2.backup_finish_date), 101), 'Not Yet Taken') AS LastBackUpTaken

FROM master.sys.databases T1 LEFT OUTER JOIN

msdb.dbo.backupset T2

ON T2.database_name = T1.name

GROUP BY T1.Name

ORDER BY T1.Name

5. To get complete backups information of a particular database

SELECT s.database_name,

m.physical_device_name,

cast(s.backup_size/1000000 as varchar(14))+ ' '+ 'MB' as bkSize,

CAST (DATEDIFF(second,s.backup_start_date , s.backup_finish_date)AS VARCHAR(4))+ ' '+ 'Seconds' TimeTaken,

s.backup_start_date,

CASE s.[type]

WHEN 'D' THEN 'Full'

WHEN 'I' THEN 'Differential'

WHEN 'L' THEN 'Transaction Log'

```

END as BackupType,

s.server_name, s.recovery_model

FROM msdb.dbo.backupset s

inner join msdb.dbo.backupmediafamily m

ON s.media_set_id = m.media_set_id

WHERE s.database_name = 'AdventureWorks'

ORDER BY database_name, backup_start_date, backup_finish_date

```

6. DMV to monitor locks

```

SELECT
t1.resource_type,t1.resource_database_id,t1.resource_associated_entity_id,t1.request_mode,t1.request_session_id,

t2.blocking_session_id,o1.name 'object name',o1.type_desc 'object

descr',p1.partition_id 'partition id',p1.rows 'partition/page rows',

a1.type_desc 'index descr',a1.container_id 'index/page container_id' FROM

sys.dm_tran_locks as t1

INNER JOIN sys.dm_os_waiting_tasks as t2

ON t1.lock_owner_address = t2.resource_address

LEFT OUTER JOIN sys.objects o1 on o1.object_id =

t1.resource_associated_entity_id

LEFT OUTER JOIN sys.partitions p1 on p1.hobt_id =

t1.resource_associated_entity_id

LEFT OUTER JOIN sys.allocation_units a1 on a1.allocation_unit_id =

t1.resource_associated_entity_id

```

7. Displaying expensive queries

```

SELECT TOP 10 SUBSTRING(qt.TEXT, (qs.statement_start_offset/2)+1,

((CASE qs.statement_end_offset

```

```

WHEN -1 THEN DATALENGTH(qt.TEXT)

ELSE qs.statement_end_offset

END - qs.statement_start_offset)/2)+1),

qs.execution_count,

qs.total_logical_reads, qs.last_logical_reads,

qs.total_logical_writes, qs.last_logical_writes,

qs.total_worker_time,

qs.last_worker_time,

qs.total_elapsed_time/1000000 total_elapsed_time_in_S,

qs.last_elapsed_time/1000000 last_elapsed_time_in_S,

qs.last_execution_time,

qp.query_plan

FROM sys.dm_exec_query_stats qs

CROSS APPLY sys.dm_exec_sql_text(qs.sql_handle) qt

CROSS APPLY sys.dm_exec_query_plan(qs.plan_handle) qp

ORDER BY qs.total_logical_reads DESC -- logical reads

-- ORDER BY qs.total_logical_writes DESC -- logical writes

-- ORDER BY qs.total_worker_time DESC -- CPU time

```

8. Expensive queries we can check using

Activity Monitor

9. Using server dashboard report

Performance - Top Queries by Total CPU

IO

10. To view cached plans

```

SELECT cp.objtype AS PlanType,

       OBJECT_NAME(st.objectid,st.dbid) AS ObjectName,

       cp.ref countscp.refcounts AS ReferenceCounts,

       cp.usecounts AS UseCounts,

       st.text AS SQLBatch,

       qp.query_plan AS QueryPlan

FROM sys.dm_exec_cached_plans AS cp

CROSS APPLY sys.dm_exec_query_plan(cp.plan_handle) AS qp

CROSS APPLY sys.dm_exec_sql_text(cp.plan_handle) AS st;

GO

```

Note: To view the query in a particular session we can use

```
dbcc inputbuffer(spids)
```

But this command displays ~~only the first~~~~only first~~ 256 chars of a query/batch.

11. To view complete query we can use the following DMF from SS 2005

```
sys.dm_exec_sql_text
```

12. To view no of cached plans in procedure cache we can use

```
dbcc proccache
```

13. To remove execution plans from procedure cache

```
dbcc freeproccache
```

14. Recovery model, log reuse wait description, log file size, log usage size and compatibility level for all databases on instance

```

SELECT db.[name] AS [Database Name],

db.recovery_model_desc AS [Recovery Model],

db.log_reuse_wait_desc AS [Log Reuse Wait Description],

ls.cntr_value AS [Log Size (KB)], lu.cntr_value AS [Log Used (KB)],

CAST(CAST(lu.cntr_value AS FLOAT) / CAST(ls.cntr_value AS FLOAT)AS DECIMAL(18,2)) * 100 AS [Log Used %],

db.[compatibility_level] AS [DB Compatibility Level], db.page_verify_option_desc AS [Page Verify Option]

FROM sys.databases AS db

INNER JOIN sys.dm_os_performance_counters AS lu

ON db.name = lu.instance_name

INNER JOIN sys.dm_os_performance_counters AS ls

ON db.name = ls.instance_name

WHERE lu.counter_name LIKE 'Log File(s) Used Size (KB)%'

AND ls.counter_name LIKE 'Log File(s) Size (KB)%';

```

15. Backup and Restoration

Note: To check recent backups of all databases

```

SELECT

T1.Name AS DatabaseName,

COALESCE(CONVERT(VARCHAR(12), MAX(T2.backup_finish_date), 101), 'Not Yet Taken') AS LastBackUpTaken

FROM master.sys.databases T1 LEFT OUTER JOIN

msdb.dbo.backupset T2

ON T2.database_name = T1.name

```

GROUP BY T1.Name

ORDER BY T1.Name

16. To get complete backups information of a particular database

```
SELECT s.database_name,  
  
m.physical_device_name,  
  
cast(s.backup_size/1000000 as varchar(14))+ ' '+ 'MB' as bkSize,  
  
CAST (DATEDIFF(second,s.backup_start_date , s.backup_finish_date)AS VARCHAR(4))+ ' '+ 'Seconds' TimeTaken,  
  
s.backup_start_date,  
  
CASE s.[type]  
  
WHEN 'D' THEN 'Full'  
  
WHEN 'I' THEN 'Differential'  
  
WHEN 'L' THEN 'Transaction Log'  
  
END as BackupType,  
  
s.server_name, s.recovery_model  
  
FROM msdb.dbo.backupset s  
  
inner join msdb.dbo.backupmediafamily m  
  
ON s.media_set_id = m.media_set_id  
  
WHERE s.database_name = 'AdventureWorks'  
  
ORDER BY database_name, backup_start_date, backup_finish_date
```

ISSUE 16: DBCC & SP'S

DBCC:

1.DBCC CHECKALLOC

DBCC CHECKALLOC checks page usage and allocation in the database. Use this command if allocation errors are found for the database. If you run DBCC CHECKDB, you do not need to run DBCC CHECKALLOC, as DBCC CHECKDB includes the same checks (and more) that DBCC CHECKALLOC performs.

2.DBCC CHECKCATALOG

This command checks for consistency in and between system tables. This command is not executed within the DBCC CHECKDB command, so running this command weekly is recommended.

3.DBCC CHECKCONSTRAINTS

DBCC CHECKCONSTRAINTS alerts you to any CHECK or constraint violations.

Use it if you suspect that there are rows in your tables that do not meet the constraint or CHECK constraint rules.

4.DBCC CHECKDB

A very important DBCC command, DBCC CHECKDB should run on your SQL Server instance on at least a weekly basis. Although each release of SQL Server reduces occurrences of integrity or allocation errors, they still do happen. DBCC CHECKDB includes the same checks as DBCC CHECKALLOC and DBCC CHECKTABLE. DBCC CHECKDB can be rough on concurrency, so be sure to run it at off-peak times.

5.DBCC CHECKTABLE

DBCC CHECKTABLE is almost identical to DBCC CHECKDB, except that it is performed at the table level, not the database level. DBCC CHECKTABLE verifies index and data page links, index sort order, page pointers, index pointers, data page integrity, and page offsets. DBCC CHECKTABLE uses schema locks by default, but can use the TABLOCK option to acquire a shared table lock. CHECKTABLE also performs object checking using parallelism by default (if on a multi-CPU system).

6.DBCC CHECKFILEGROUP

DBCC CHECKFILEGROUP works just like DBCC CHECKDB, only DBCC CHECKFILEGROUP checks the specified filegroup for allocation and structural issues. If you have a very large database (this term is relative, and higher end systems may be more apt at performing well with multi-GB or TB systems) , running DBCC CHECKDB may be time-prohibitive.

If your database is divided into user defined filegroups, DBCC CHECKFILEGROUP will allow you to isolate your integrity checks, as well as stagger them over time.

7.DBCC CHECKIDENT

DBCC CHECKIDENT returns the current identity value for the specified table, and allows you to correct the identity value if necessary.

8.DBCC DBREINDEX

If your database allows modifications and has indexes, you should rebuild your indexes on a regular basis. The frequency of your index rebuilds depends on the level of database activity, and how quickly your database and indexes become fragmented. DBCC DBREINDEX allows you to rebuild one or all indexes for a table. Like DBCC CHECKDB, DBCC CHECKTABLE, DBCC CHECKALLOC, running DBREINDEX during peak activity times can significantly reduce concurrency.

9.DBCC INDEXDEFRAG

Microsoft introduced the excellent DBCC INDEXDEFRAG statement beginning with SQL Server 2000. This DBCC command, unlike DBCC DBREINDEX, does not hold long term locks on indexes. Use DBCC INDEXDEFRAG for indexes that are not very fragmented, otherwise the time this operation takes will be far longer than running DBCC DBREINDEX. In spite of ~~its~~^{its} ability to run during peak periods, DBCC INDEXDEFRAG has had limited effectiveness compared to DBCC DBREINDEX (or drop/create index).

10.DBCC INPUTBUFFER

The DBCC INPUTBUFFER command is used to view the last statement sent by the client connection to SQL Server. When calling this DBCC command, you designate the SPID to examine. (SPID is the process ID, which you can get from viewing current activity in Enterprise Manager or executing sp_who.)

11.DBCC OPENTRAN

DBCC OPENTRAN is a Transact-SQL command that is used to view the oldest running transaction for the selected database. The DBCC command is very useful for troubleshooting orphaned connections (connections still open on the database but disconnected from the application or client), and identification of transactions missing a COMMIT or ROLLBACK. This command also returns the oldest distributed and undistributed replicated transactions, if any exist within the database. If there are no active transactions, no data will be returned. If you are having issues with your

transaction log not truncating inactive portions, DBCC OPENTRAN can show if an open transaction may be causing it.

12.DBCC PROCCACHE

You may not use this too frequently, however it is an interesting DBCC command to execute periodically, particularly when you suspect you have memory issues. DBCC PROCCACHE provides information about the size and usage of the SQL Server procedure cache.

13.DBCC SHOWCONTIG

The DBCC SHOWCONTIG command reveals the level of fragmentation for a specific table and its indices. This DBCC command is critical to determining if your table or index has internal or external fragmentation. Internal fragmentation concerns how full an 8K page is.

When a page is underutilized, more I/O operations may be necessary to fulfill a query request than if the page was full, or almost full.

External fragmentation concerns how contiguous the extents are. There are eight 8K pages per extent, making each extent 64K. Several extents can make up the data of a table or index. If the extents are not physically close to each other, and are not in order, performance could diminish.

14.DBCC SHRINKDATABASE

DBCC SHRINKDATABASE shrinks the data and log files in your database.

Avoid executing this command during busy periods in production, as it has a negative impact on I/O and user concurrency. Also remember that you cannot shrink a database past the target percentage specified, shrink smaller than the model database, shrink a file past the original file creation size, or shrink a file size used in an ALTER DATABASE statement.

15.DBCC SHRINKFILE

DBCC SHRINKFILE allows you to shrink the size of individual data and log files. (Use sp_helpfile to gather database file ids and sizes).

16. DBCC TRACEOFF, TRACEON, TRACESTATUS

Trace flags are used within SQL Server to temporarily enable or disable specific SQL Server instance characteristics.

Traces are enabled using the DBCC TRACEON command, and disabled using DBCC TRACEOFF. DBCC TRACESTATUS is used to ~~display~~~~displays~~ the status of trace flags. You'll most often see TRACEON used in conjunction with deadlock logging (providing more verbose error information).

17.DBCC USEROPTIONS

Execute DBCC USEROPTIONS to see what user options are in effect for your specific user connection. This can be helpful if you are trying to determine if ~~your~~~~you~~ current user options are inconsistent with the database options.

18. DBCC SQLPERF(LOGSPACE) - To check the current size of log(.LDF) files of all the databases.

(in case of disk space issue or on log file autogrowth error)

19.DBCC OPENTRAN - To check the active transaction(s) of the current database.

20. DBCC ERRORLOG: If you rarely restart SQL Server service, resulting server log gets very large and takes a long time to load and view. You can truncate (essentially create a new log) the Current Server log by this.

You can accomplish the same thing using this stored procedure: sp_cycle_errorlog.

21. DBCC DROPCLEANBUFFERS: To remove all the data from SQL Server's data cache (buffer) between performance tests to ensure fair testing. Fyi, this command only removes clean buffers, not dirty buffers.

So, before running the DBCC DROPCLEANBUFFERS command, you may first want to run the CHECKPOINT command.

Running CHECKPOINT will write all dirty buffers to disk. So, when you run DBCC DROPCLEANBUFFERS, you can be assured that all data buffers are cleaned out, not just the clean ones.

22. DBCC updatestatistics

GENERAL HELP PROCEDURES:

<u>sp_depends</u>	Better version of sp_depends
<u>sp_help</u>	Better sp_help
<u>sp_helpdb</u>	Database Information
<u>sp_helpdevice</u>	Break down database devices into a nice report
<u>sp_helpgroup</u>	List groups in database by access level
<u>sp_helpindex</u>	Shows indexes by table
<u>sp_helpsegment</u>	Segment Information
<u>sp_helpprotect</u>	Simple Protection Info for the database
<u>sp_helptext</u>	Show comments with line splits ok
<u>sp_helpuser</u>	Lists users in current database by group (includes aliases)
<u>sp_lock</u>	Lock information
<u>sp_syntax</u>	Works on any procedure to give you syntax
<u>sp_who</u>	sp_who that fits on a page

SYSTEM ADMINISTRATOR PROCEDURES:

<u>sp_block</u>	Blocking processes.
<u>sp_dbSPACE</u>	Summary of current database space information.
<u>sp_dumpdevice</u>	Listing of Dump devices
<u>sp_diskdevice</u>	Listing of Disk devices
<u>sp_helpdbdev</u>	Show how Databases use Devices
<u>sp_helplogin</u>	Show logins and remote logins to server

<u>sp_helpmirror</u>	Shows mirror information, discover broken mirrors
<u>sp_segment</u>	Segment Information
<u>sp_server</u>	Server summary report (very useful)
<u>sp_stat</u>	Give basic server performance information (loops)
<u>sp_vdevno</u>	Who's who in the device world

DBA PROCEDURES:

<u>sp_badindex</u>	list badly formed indexes (allow nulls) or those needing statistics
<u>sp_collist</u>	list all columns in database
<u>sp_find_missing_index</u>	Finds keys that do not have associated index
<u>sp_flowchart</u>	Makes a flowchart of procedure nesting
<u>sp_groupprotect</u>	Permission info by group
<u>sp_indexspace</u>	Space used by indexes in database
<u>sp_id</u>	Gives information on who you are and which db you are in
<u>sp_noindex</u>	list of tables without indexes.
<u>sp_helpcolumn</u>	show columns for given table
<u>sp_helpdefault</u>	list defaults (part of object list objectlist)
<u>sp_helpobject</u>	list objects
<u>sp_helpproc</u>	list procs (part of object list objectlist)
<u>sp_helprule</u>	list rules (part of object list objectlist)
<u>sp_helptable</u>	list tables (part of object list objectlist)
<u>sp_helptrigger</u>	list triggers (part of object list objectlist)
<u>sp_helpview</u>	list views (part of object list objectlist)
<u>sp_objprotect</u>	Permission info by object
<u>sp_read_write</u>	list tables by # procs that read, # that write, # that do both
<u>sp_trigger</u>	Useful synopsis report of current database trigger schema
<u>sp_whodo</u>	sp__who - filtered for only active processes

AUDIT PROCEDURES:

[sp_auditsecurity](#) Security Audit On Server

[sp_auditdb](#) Audit Current Database For Potential Problems

[sp_checkkey](#) Generate script for referential integrity problems (uses key info from sp_foreignkey)

REVERSE ENGINEERING PROCEDURES:

[sp_revalias](#) get alias generation script for current database

[sp_revdb](#) get database generation script for server

[sp_revdevice](#) get device generation script for server

[sp_revgroup](#) get group generation script for current database

[sp_revindex](#) get index generation script for current database

[sp_revlogin](#) get login generation script for server

[sp_revmirror](#) get mirror generation script for current database

[sp_revsegment](#) get segment generation script for current database

[sp_revtable](#) get table generation script for current database

[sp_revuser](#) get user generation script for current database

OTHER PROCEDURES:

[sp_bcp](#) Create unix script to bcp in/out database

[sp_date](#) Who can remember all the date styles?

[sp_iostat](#) Loops n times showing active processes only

[sp_grep](#) Search for ~~pattern~~ **pattern**

[sp_isactive](#) Shows info about a single active process

[sp_ls](#) Lists specific objects

[sp_quickstats](#) Quick dump of server summary information

[sp_whoactive](#) Show info about who is active

ISSUE 17: QUERY ARCHITECTURE

Performance Tuning, Monitoring and Troubleshooting

* As part of performance tuning we have to analyze and work with

* Physical I/O and Logical I/O

- * CPU usage
- * Memory usage
- * Database Design
- * Application's db programming methods

Query Architecture

- * Once the query is submitted to Database Engine for first time it performs the following tasks.
 - * Parsing (Compiling)
 - * Resolving (Verifying syntax, table, col names etc)
 - * Optimizing (Generating execution plan)
 - * Executing (Executing query)
- * For next time if the query was executed with same case and same no of characters i.e with no extra spaces then the query is executed by taking existing plan.
- * To display cached plans


```
SELECT cp.objtype AS PlanType,
       OBJECT_NAME(st.objectid,st.dbid) AS ObjectName,
       cp.refcounts AS Reference Counts, cp. use countsusecounts AS UseCounts,
       st.text AS SQLBatch,qp.query_plan AS QueryPlan
FROM sys.dm_exec_cached_plans AS cp
CROSS APPLY sys.dm_exec_query_plan(cp.plan_handle) AS qp
CROSS APPLY sys.dm_exec_sql_text(cp.plan_handle) AS st;
GO
```
- * To remove plans from cache memory


```
DBCC FREEPROCCACHE
```

Execution Plan

- * Step by step process followed by SS to execute a query is ~~called an execution~~called execution plan.
- * It is prepared by Query Optimizer using STATISTICS.
- * Query optimizer ~~prepares the execution~~prepares execution plan and ~~stores it in~~stores in the ProcedureCache.
- * Execution plans are different for
 - * Different case statements
 - * Different size statements (spaces.)
- * To view graphical execution plan
 - * select the query --> press ctrl+M/L
- * To view xml execution plan
 - * set showplan_xml on/off
 - * Execute the query
- * To view text based execution plan
 - * set showplan_text on/off
 - * Execute the query.

Statistics

- * Consists of meta data of the table or index.
- * If statistics are out of date, ~~the query~~query optimizer may ~~prepare a poor~~prepare poor plan.
- * We have to update statistics weekly ~~with a maintenance~~with maintenance plan.

USE master

GO

-- Enable Auto Update of Statistics

ALTER DATABASE AdventureWorks SET AUTO_UPDATE_STATISTICS ON;

GO

-- Update Statistics for whole database

EXEC sp_updatestats

GO

-- Get List of All the Statistics of Employee table

```

sp_helpstats 'Human Resources .Employee', 'ALL'
GO
-- Get List of statistics of AK_Employee_NationalIDNumber index
DBCC SHOW_STATISTICS ("HumanResources.Employee",AK_Employee_NationalIDNumber)
-- Update Statistics for single table
STATISTICSUPDATE STATISTICS Human Resources. Employee
GO
-- Update Statistics for single index on single table
UPDATE STATISTICS Human Resources.Employee AK_Employee_NationalIDNumber
GO

```

Index

- * It is another database objects which can be used
 - * To reduce searching process
 - * To enforce uniqueness
- * By default SS ~~searches~~search for the rows by following the process called table scan.
- * If the table consists of huge data then table scan provides less performance.
- * Index is created ~~in a tree-like~~in tree-like structure which consists of root, node and leaf level.
- * At leaf level, index pages are present by default.
- * We can place max 250 indexes per table.
- * Indexes are automatically placed if we place
 - * Primary key (clustered)
 - * Unique (unique non clustered index)
- * We can place indexes as follows


```
create [unique][clustered/nonclustered] index <indexName>
on <name><name>/<viewName>(col1,col2,...)
[include(...)]
```

Types

- * Clustered
- * NonClustered

1. Clustered Index-----

- * It physically sorts the rows in the table.
- * A table can have only ONE clustered index.
- * Both data and index pages are merged and stored at third level (Leaf level).
- * We can place on columns which are used to search a range of rows,

Ex:

```

Create table prods(pid int,name varchar(40), qty int)
insert prods values(4,'Books',50),(2,'Pens',400)

```

```
select * from prods (run the query by pressing ctrl+L)
```

```
create clustered index pid_idx on prods(pid)
```

```
select * from prods -- check the rows are sorted in asc order to pid
```

FAQ:- Difference between P.K and Clustered Index?

* Primary ~~keys~~key enforce uniqueness and ~~allow establishing~~allows to ~~establish~~establish relationship. But by default clustered index cannot.

```
select * from prods where pid=2 -- press ctrl+L to check execution plan
```

```
insert prods values(3,'Pencils',500) -- Check this row is inserted as a secondas second record.
```

Note: A table without clustered index is called HEAP where the rows and pages of the table are not present in any order.

NonClustered Index-----

- * It cannot sort the rows physically.
- * We can place max 249 nonclustered indexes on the table.
- * Both data and index pages are stored separately.
- * It locates rows either from heap (Table scan) or from clustered index.
- * Always we have to place first clustered index then nonclustered index.
- * If the table is heap the index page consists of
IndexKeyColvalues row reference
- * If the table consists of clustered index then index page consists of
IndexKeyColValues Clustered Index Keycol Values
- * Nonclustered indexes are rebuilt when
 - * Clustered index is created/dropped/modified

Ex: Create a nonclustered index on the name column of the periodic prods table.

```
create index index1 on prods(pname)
select * from prods where pname='Books' -- check execution plan
```

- * To disp indexes present on a table
sp_helpindex <tname>
- * To drop index
drop index prods.pid_indx
- * To disp space used by the index
sp_spaceused prods

Using Included Columns in NonClustered Index-----

- * We can maintain regularly used columns in nonclustered index so that no need that SQL Server should take data from heap or clustered index.
- * If the number of rows is more it provides better performance.

Ex:

```
--step1
USE AdventureWorks
GO
CREATE NONCLUSTERED INDEX IX_Address_PostalCode
ON Person.Address (PostalCode)
INCLUDE (AddressLine1, AddressLine2, City, StateProvinceID)
GO
```

```
--step2
SELECT AddressLine1, AddressLine2, City, StateProvinceID, PostalCode
FROM Person.Address
WHERE PostalCode BETWEEN '98000'
AND '99999';
GO
```

Index Management

Fill Factor-----

- * Percentage of space used in leaf level index pages.
- * By default it is 100%.

- * To reduce page splits when the data is manipulated in the base table we can set proper FillFactor.
- * It allows online index processing
 - * While the index rebuilding process is going on users can work with the table.

Page Split-----

- * Due to regular changes in the table if the index pages are full to allocate memory for the index key columns SS takes remaining rows into new page. This process is called Page split.
- * Page split increases size of index and the index pages order changes.
- * This situation where unused free space is available and the index pages are not in the order of key column values is called fragmentation.
- * To find fragmentation level we can use
 - dbcc showcontig
 - or
 - We can use sys.dm_db_index_physical_stats DMF as follows

```
SELECT a.index_id, name, avg_fragmentation_in_percent
FROM sys.dm_db_index_physical_stats
(DB_ID('AdventureWorks'),
OBJECT_ID('Production.Product'), NULL, NULL, NULL)
AS a JOIN sys.indexes AS b
ON a.object_id = b.object_id
AND a.index_id = b.index_id;
```

- * To control fragmentation we can either reorganize the index or rebuild the index.

1. Reorganizing Index * It is the process of arranging the index pages according to the order of index key column values.

- * If the fragmentation level is more than 5 to 8% and less than ~~28 to 28%~~ 30% then we can reorganize the indexes.
- * It cannot reduce the index ~~size as well as~~ statistics are not updated.

syn:

```
ALTER INDEX <indexName>/<All> on <tname> REORGANIZE
```

2. Index Rebuilding * It is the process of deleting and creating fresh ~~indexes~~index.

- * It reduces the size of index and updates statistics
- * If the fragmentation level is more than 30% then we can rebuild indexes.

syn:

```
ALTER INDEX <indexName>/<ALL> on <tname> REBUILD
```

Note:

If we have ~~mentioned the ONLINE~~mentioned-ONLINE INDEX PROCESSING option then rebuilding takes space in TEMPDB.

To check ~~consistency~~consistancy of a database we can use DBCC CHECKDB('dbName') ~~if it disp if~~ any corrupted pages are present, use space in tempdb.

Transactions and Locks

* A transaction ~~is a single~~is-single unit of work which may ~~consiste~~consists of one or more commands.

* Transactions works with ACID properties

- * ~~Atomicity~~Atomicity
- * ~~Consistency~~Consistancy
- * Isolation
- * Durability

* SQL Server supports 2 types of transactions

- * Implicit
- * Explicit

* By default SS supports implicit transaction where for every insert, update and delete 3 records are stored in T.Log file

Begin tran

insert/update/delete

commit tran

* To implement business logic i.e. according to the required if we want to commit or can use explicit transactions.

rollback the changes we

Begin Tran

commit/rollback tran

- * Any transaction which consists of manipulations places locks on the tables.
- * By default when we make a db as current db automatically Shared Lock is placed.
- * While working with insert,update,delete by default SS places Exclusive lock.
- * Type of locks placed on objects depends on isolation levels.

Isolation Levels

- * It is a transaction property.
- * Types of locks placed by SS on the resource ~~depend~~depends on isolation levels.
- * SS supports 5 isolation levels
 - * Read Committed (Default)
 - * Read Uncommitted
 - * Repeatable Reads
 - * Snapshot
 - * Serializable
- * To check the isolation level
dbcc useroptions
- * To set the isolation level
SET TRANSACTION ISOLATION LEVEL <~~required isolationlevel~~requiredisolationlevel>
- * To handle the concurrency related problems SS places locks
- * SS supports 2 types of concurrencies
 - * Optimistic Concurrency
 - * Uses Shared Locks
 - * More concurrency
 - * Pessimistic Concurrency
 - * Uses Exclusive Locks
 - * Low concurrency

Ex: Open new query window

--user1

use Test

go

begin tran

update emp set sal=5000

Take new query -->

--user2

use Test

go

select * from emp (--query runs ~~continuously~~continuesly till user1 session releases lock)

Take new query

--user3

set transaction isolation level read uncommitted

select * from emp

--Take new query

sp_lock -- To view locks information

```
or
select * from sys.dm_tran_locks

--check blocking using
sp_who/sp_who2

-- To check locks placed by a particular session
sp_lock <spid>
sp_lock 56
```

ISSUE 18: PROFILER EVENTS

What Data to Collect:

Profiler allows you to specify which events you want to capture and which data columns from those ~~event~~event to capture. In addition, you can use filters to reduce the incoming data to only what you need for this specific analysis.

Events to Capture:

- ? Stored Procedures--RPC:Completed
- ? TSQL--SQL:BatchCompleted

You may be surprised that only two different events need to be captured: one for capturing stored procedures and one for capturing all other Transact-SQL queries.

Data Columns to Capture:

- ? Duration (data needs to be grouped by duration)
- ? Event Class
- ? DatabaseID (If you have more than one database on the server)
- ? TextData
- ? CPU
- ? Writes
- ? Reads
- ? StartTime (optional)
- ? EndTime (optional)
- ? ApplicationName (optional)
- ? NTUserName (optional)
- ? LoginName (optional)

SPID

The data you want to actually capture and view includes some that are very important to you, especially duration and TextData; and some that are not so important, but can be useful, such as ApplicationName or NTUserName.

Filters to Use:

- Duration > 5000 milliseconds (5 seconds)
- Don't collect system events
- Collect data by individual database ID, not all databases at once
- Others, as appropriate

Filters are used to reduce the amount of data collected, and the more filters you use, the more data you can filter out that is not important. Generally, I use three filters, but others can be used, as appropriate to your situation. And of these, the most important is duration. I only want to collect information on those that have enough duration to be of importance to me, as we have already discussed.

Collecting the Data:

Depending on the filters you used, and the amount of time you run Profiler to collect the data, and how busy your server is, you may collect a lot of rows of data. While you have several choices, I suggest you configure Profiler to save the data to a file on ~~your~~ your local computer (not on the server you are Profiling), and not set a maximum file size. Instead, let the file grow as big as it needs to grow. You may want to watch the growth of this file, in case it gets out of hand. In most cases, if you have used appropriate filters, the size should stay manageable. I recommend using one large file because it is easier to identify long running queries if you do.

As mentioned before, collect your trace file during a typical production period, over a period of 3-4 hours or so. As the data is being collected, it will be sorted for you by duration, with the longest running queries appearing at the bottom of the Profiler window. It can be interesting to watch this window for ~~a while~~ awhile while you are collecting data. If you like, you can configure Profiler to automatically turn itself off at the appropriate time, or you can do this manually.

Once the time is up and the trace stopped, the Profiler trace is now stored in the memory of the local computer, and on disk. Now you are ready to identify those long running queries.

Analyzing the Data:

Guess what, you have already identified all queries that ran during the trace collection that exceed your specified duration, whatever it was. So if you selected a duration of 5 seconds, you will only see those queries that took longer than five seconds to run. By definition, all the queries you have captured need to be tuned. "What! But over 500 queries were captured! That's a lot of work!" It is not as bad as you think. In most cases, many of the queries you have captured are duplicate queries. In other words, you have probably captured the same query over and over again in your trace. So those 500 captured queries may only be 10, or 50, or even 100 distinct queries. On the other hand, there may be only a handful of queries captured (if you are lucky).

Whether you have just a handful, or a lot of slow running queries, ~~your~~^{you} next job is to determine which are the most critical for you to analyze and tune first. This is where you need to set priorities, as you probably don't have enough time to analyze them all.

To prioritize the long running queries, you will probably want to first focus on those that run the longest. But as you do this, keep in mind how often each query is run.

For example, if you know that a particular query is for a report that only runs once a month (and you happened to have captured it when it was running), and this query took 60 second to run, it probably is not as high as a priority to tune as a query that takes 10 seconds to run, but runs 10 times a minute. In other words, you need to balance the length of how long a query takes to run, to how often it runs. With this in mind, you need to identify and prioritize those queries that take the most physical SQL Server resources to run. Once you have done this, then you are ready to analyze and tune them.

Traces that you want to replay must contain a minimum set of events and data columns. If the trace doesn't contain the necessary elements, you won't be able to replay the trace. The required elements are in addition to any other elements that you want to monitor or display with traces. Events that you must capture in order to allow a trace to be replayed and analyzed correctly are

- ❓ **Connect**
- ❓ **CursorExecute** (required only when replaying server-side cursors)
- ❓ **CursorOpen** (required only when replaying server-side cursors)
- ❓ **CursorPrepare** (required only when replaying server-side cursors)
- ❓ **Disconnect**
- ❓ **Exec Prepared SQL** (required only when replaying server-side prepared SQL statements)
- ❓ **ExistingConnection**
- ❓ **Prepare SQL** (required only when replaying server-side prepared SQL statements)
- ❓ **RPC:OutputParameter**
- ❓ **RPC:Starting**
- ❓ **SQL:BatchStarting**

Data columns that you must capture to allow a trace to be replayed and analyzed correctly are:

- ❓ **Application Name**
- ❓ **Binary Data**
- ❓ **Connection ID or SPID**
- ❓ **Database ID**

?	Event Class
?	Event SubClass
?	Host Name
?	Integer Data
?	Server Name
?	SQL User Name
?	Start Time
?	Text

ISSUE 19: DMV

`sys.dm_os_wait_stats` is the DMV that contains wait statistics, which are aggregated across all session ids since the last restart of SQL Server or since the last time that the wait statistics were reset manually using `DBCC SQLPERF ('sys.dm_os_wait_stats', CLEAR)`. Resetting wait statistics can be helpful before running a test or workload.

Anytime a session_id waits for a resource, the session_id is moved to the waiter list along with an associated wait type. The DMV `sys.dm_os_waiting_tasks` shows the waiter list at a given moment in time. Waits for all session_ids are aggregated in `sys.dm_os_wait_stats`.

The stored procedures `track_waitstats_2005` and `get_waitstats_2005` can be used to measure the wait statistics for a given workload.

What are DMVs

Dynamic Management Views are views and functions introduced in sql server 2005 for monitoring and tuning sql server performance Two types of dynamic management views:

- Server-scoped DMV: Stored in Master Database
- Database-scoped DMV: Specific to each database

Permission to Execute DMV [Security]

To query a server scoped DMV, the database user must have **SELECT** privilege on **VIEW SERVER STATE** and for database scoped DMV, the user must have **SELECT** privilege on **VIEW DATABASE STATE**.

- GRANT VIEW SERVER STATE to <Login>
- GRANT VIEW DATABASE STATE to <User>

If you want to deny a user permission to query certain DMVs, you can use the **DENY** command to restrict access to a specific DMV.

All the DMVs ~~exist~~^{exists} in SYS schema and their names start with **DM_**. So when you need to query a DMV, you should prefix the view name with **SYS**. As an example, if you need to see the total physical memory of the SQL Server machine;

```
SELECT  
(Physical_memory_in_bytes/1024.0)/1024.0 AS Physical_memory_in_Mb  
FROM sys.dm_os_sys_info
```

how many DMV/DMF are there in SQL Server, to get that information (see Pinal's [post](#))

```
SELECT name, type, type_desc FROM sys.system_objects WHERE name LIKE 'dm_%' ORDER BY name
```

or

```
SELECT name, type, type_desc FROM sys.system_objects WHERE name LIKE 'dm[_]%' ORDER BY name
```

Frequently used

1. SQL Server related [Hardware Resources] DMV
2. Database related DMV
3. Index related DMV
4. Execution related DMV

1. SQL Server Related DMV

This section details the DMVs associated with SQL Server ~~system~~^{system}. SQL DMV is responsible ~~for managing to~~^{for managing} server level resources specific to a SQL Server instance.

This section covers DMVs related to OS, Disk and Memory.

a. sys.dm_os_sys_info

This view returns the information about the SQL Server machine, available resources and the resource consumption.

This view returns information like the following:

1. CPU Count: Number of logical CPUs in the server
2. Hyperthread-ratio: Ratio of logical and physical CPUs
3. **Physical_memory_in_bytes**: Amount of physical memory available
4. **Virtual_memory_in_bytes**: Amount of virtual memory available
5. **Bpool_committed**: Committed physical memory in buffer pool
6. **OS_Priority_class**: Priority class for SQL Server process
7. **Max_workers_thread**: Maximum number of workers which can be created

b. sys.dm_os_hosts

This view returns all the hosts registered with SQL Server 2005. This view also provides the resources used by each host.

1. **Name:** Name of the host registered
2. **Type:** Type of hosted component [SQL Native Interface/OLE DB/MSDART]
3. **Active_tasks_count:** Number active tasks host placed
4. **Active_ios_count:** I/O requests from host waiting

c. sys.dm_os_schedulers

sys.dm_os_schedulers view will help you identify if there is any CPU bottleneck in the SQL Server machine. The number of runnable tasks is generally a nonzero value; a nonzero value indicates that tasks have to wait for their time slice to run. If the runnable task counts show high values, then there is a symptom of CPU bottleneck.

 [Collapse](#) | [Copy Code](#)

```
SELECT
scheduler_id,current_tasks_count,runnable_tasks_count
FROM sys.dm_os_schedulers
WHERE scheduler_id < 255
```

The above query will list all the available schedulers in the SQL Server machine and the number of runnable tasks for each scheduler.

d. sys.dm_io_pending_io_requests

This dynamic view will return the I/O requests pending ~~on the SQL~~**in SQL** Server side. It gives you information like:

1. **Io_type:** Type of pending I/O request
2. **Io_pending:** Indicates whether the I/O request is pending or has been completed by Windows
3. **Scheduler_address:** Scheduler on which this I/O request was issued

e. sys.dm_io_virtual_file_stats

This view returns I/O statistics for data and log files [MDF and LDF file]. This view is one of the commonly used views and will help you to identify I/O file level. This will return information like:

1. **Sample_ms:** Number of milliseconds since the instance of SQL Server has started
2. **Num_of_reads:** Number of reads issued on the file
3. **Num_of_bytes_read:** Total number of bytes read on this file
4. **Io_stall_read_ms:** Total time, in milliseconds, that the users waited for reads issued on the file
5. **Num_of_writes:** Number of writes made on this file
6. **Num_of_bytes_written:** Total number of bytes written to the file
7. **Io_stall_write_ms:** Total time, in milliseconds, that users waited for writes to be completed on the file
8. **Io_stall:** Total time, in milliseconds, that users waited for I/O to be completed
9. **Size_on_disk_bytes:** Number of bytes used on the disk for this file

f. sys.dm_os_memory_clerks

This DMV will help how much memory SQL Server has allocated through AWE.

 [Collapse](#) | [Copy Code](#)

```
SELECT
SUM(awe_allocated_kb) / 1024 as [AWE allocated, Mb]
FROM sys.dm_os_memory_clerks
```

The same DMV can be used to get the memory consumption by internal components of SQL Server 2005.

 [Collapse](#) | [Copy Code](#)

```
SELECT TOP 10 type,
SUM(single_pages_kb) as [SPA Mem, Kb]
FROM sys.dm_os_memory_clerks
GROUP BY type
ORDER BY SUM(single_pages_kb) DESC
```

g. sys.dm_os_ring_buffers

This DMV uses **RING_BUFFER_RESOURCE_MONITOR** and gives information from resource monitor notifications to identify memory state changes. Internally, SQL Server has a framework that monitors different memory pressures. When the memory state changes, the resource monitor task generates a notification. This notification is used internally by the components to adjust their memory usage according to the memory state.

 [Collapse](#) | [Copy Code](#)

```
SELECT
Record FROM sys.dm_os_ring_buffers
WHERE ring_buffer_type = 'RING_BUFFER_RESOURCE_MONITOR'
```

The output of the above query will be in XML format. The output will help you in detecting any low memory notification.

RING_BUFFER_OOM: Ring buffer oom contains records indicating server out-of-memory conditions.

 [Collapse](#) | [Copy Code](#)

```
SELECT
record FROM sys.dm_os_ring_buffers
WHERE ring_buffer_type = 'RING_BUFFER_OOM'
```

2. Database Related DMV

This section details the DMVs associated with SQL Server Databases. These DMVs will help to identify database space usages, partition usages, session information usages, etc...

a. sys.dm_db_file_space_usage

This DMV provides the space usage information ~~of the TEMPDB~~ **of TEMPDB** database.

b. sys.dm_db_session_space_usage

This DMV provides the number of pages allocated and de-allocated by each session for the database

c. sys.dm_db_partition_stats

This DMV provides page and row-count information for every partition in the current database.

The below query shows all counts for all partitions of all indexes and heaps in the MSDB database:

 [Collapse](#) | [Copy Code](#)

```
USE MSDB;
GO
SELECT * FROM sys.dm_db_partition_stats;
```

The following query shows all counts for all partitions of Backup set table and its indexes

 [Collapse](#) | [Copy Code](#)

```
USE MSDB
GO
SELECT * FROM sys.dm_db_partition_stats
WHERE object_id = OBJECT_ID('backupset');
```

d. sys.dm_os_performance_counters

Returns the SQL Server / Database related counters maintained by the server.

The below sample query uses the **dm_os_performance_counters** DMV to get the Log file usage for all databases in KB.

 [Collapse](#) | [Copy Code](#)

```
SELECT instance_name
,cntr_value 'Log File(s) Used Size (KB)'
FROM sys.dm_os_performance_counters
WHERE counter_name = 'Log File(s) Used Size (KB)'
```

3. INDEX Related DMV

This section details the DMVs associated with SQL Server Databases. These DMVs will help to identify database space usages, Partition usages, Session information usages, etc.

a. sys.dm_db_index_usage_stats

This DMV is used to get useful information about the index usage for all objects in all databases. This also shows the amount of seeks and ~~scans~~scans for each index.

 [Collapse](#) | [Copy Code](#)

```
SELECT object_id, index_id, user_seeks, user_scans, user_lookups
```

```
FROM sys.dm_db_index_usage_stats
ORDER BY object_id, index_id
```

All indexes which have not been used so far in as database can be identified using the below Query:

 [Collapse](#) | [Copy Code](#)

```
SELECT object_name(i.object_id),
i.name,
s.user_updates,
s.user_seeks,
s.user_scans,
s.user_lookups
from sys.indexes i
left join sys.dm_db_index_usage_stats s
on s.object_id = i.object_id and
i.index_id = s.index_id and s.database_id = 5
where objectproperty(i.object_id, 'IsIndexable') = 1 and
s.index_id is null or
(s.user_updates > 0 and s.user_seeks = 0
and s.user_scans = 0 and s.user_lookups = 0)
order by object_name(i.object_id)
```

Replace the **Database_id** with the database you are looking at.

4. Execution Related DMV

Execution related DMVs will provide information regarding sessions, connections, and various requests which are coming into the SQL Server.

a. sys.dm_exec_sessions

This DMV will give information on each session connected to SQL Server. This DMV is similar to running **sp_who2** or querying Master..**sysprocesses** table.

 [Collapse](#) | [Copy Code](#)

```
SELECT
session_id,login_name,
last_request_end_time,cpu_time
FROM sys.dm_exec_sessions
WHERE session_id >= 51 – All user Sessions
```

b. sys.dm_exec_connections

This DMV shows all the ~~connection~~ **connections** to SQL Server. The below query uses **sys.dm_exec_connections** DMV to get connection information. This view returns one row for each user connection (**Sessionid >=51**).

 [Collapse](#) | [Copy Code](#)

```
SELECT
```

```
connection_id,  
session_id,client_net_address,  
auth_scheme  
FROM sys.dm_exec_connections
```

c. sys.dm_exec_requests

This DMV will give details on what each connection is actually performing in SQL Server.

 [Collapse](#) | [Copy Code](#)

```
SELECT  
session_id,status,  
command,sql_handle,database_id  
FROM sys.dm_exec_requests  
WHERE session_id >= 51
```

d. sys.dm_exec_sql_text

This dynamic management function returns the text of a SQL statement given a SQL handle.

 [Collapse](#) | [Copy Code](#)

```
SELECT  
st.text  
FROM  
sys.dm_exec_requests r  
CROSS APPLY  
sys.dm_exec_sql_text(sql_handle) AS st  
WHERE r.session_id = 51
```

Conclusion

Dynamic Management views (DMV) and Dynamic Management Functions (DMF) in SQL Server 2005 give a transparent view of what is going on inside various areas of SQL Server. By using them, we will be able to query the system for information about its current state in a much more effective manner and provide solutions much faster. DMVs can be used to ~~perform~~~~performance~~ tune and for troubleshooting ~~servers~~~~server~~ and queries. This article has shown an overview of what they are and how we can use them.

Diagnosing problems in SQL Server 2000 has always been a point of concern from both developers and ~~DBAs~~~~DBA's~~. More often than not we would have had a need to use undocumented and DBCC commands which are sometimes very difficult to understand too. SQL Server 2005 on the contrary is like ~~an open~~~~open~~ book, no need to use bit based operations and undocumented column values. Welcome the introduction of Dynamic Management Views and ~~Functions~~~~Fuctions~~ a.k.a ~~DMV and~~~~DMV's and DMF~~~~DMF's~~.

From the basic definition these dynamic management views and functions very much replace all the DBCC command outputs and the pseudo table outputs. Hence it is ~~far easier~~~~far more easier~~ to detect the health of SQL Server using these

views and functions. All these are defined in the sys schema. There are two scope for these views and ~~functions~~**function**: *Server scoped* and *Database scoped*. Incidentally unlike in SQL Server 2000 now to view these objects the user needs to have SELECT permissions and VIEW SERVER/DATABASE STATE permissions. Now that I ~~mentioned SQL~~**mentioned about SQL** Server 2000, try this yourself, create a ~~read only~~**readonly** user in a database and select the sysobjects table and check the results returned in SQL Server 2000 and SQL Server 2005.

There are multiple categories in which these views and functions have been organized. The below table shows the split:

Categories	Count
dm_broker*	4
dm_clr*	4
dm_db*	12
dm_exec*	14
dm_fts*	5
dm_io*	4
dm_os*	27
dm_qn*	1
dm_repl*	4
dm_tran*	10

So we have 85 of these views and ~~functions~~**function**. To give a further split, 76 of these are views and 9 of them are functions. So these information can be queried from the system_objects system catalog table. A typical query I used was:

```
select * from sys.system_objects
Where name like 'dm_%' Order by 1
```

Each of these views and functions have different parameters or output columns and in the next couple of queries we will try to find out how to get these values.

```
-- Getting the column details of the DMV's
Select o.name, c.name, t.name, c.column_id, c.max_length, c.precision, c.scale
FROM sys.system_columns c
INNER JOIN sys.system_objects o
ON c.object_id = o.object_id
INNER JOIN sys.types t
ON c.user_type_id = t.user_type_id
```

```
Where o.name = 'dm_os_loaded_modules'
order by 1
```

In the above query we query we get the output columns for the DMV (dm_os_loaded_modules) using the system objects. In the above query we get details like name of the output column, datatype and other length specific values. Even though this will not get us the values for the table valued functions. We will have to tweak the above query for DMF's.

```
-- Getting the column details of the DMFDMF's
Select o.name, t.name, p.*
FROM sys.system_parameters p
INNER JOIN sys.system_objects o
ON p.object_id = o.object_id
INNER JOIN sys.types t
ON p.user_type_id = t.user_type_id
Where o.name = 'dm_exec_sql_text'
order by 1
```

In the above query we try to get the parameters for the DMF (dm_exec_sql_text) using the system_parameters system catalog. So the output would show the above DMF has a parameter @handle. So if we queried this function for the sql text for a given query in the cache. The handle can be ~~obtained~~get from dm_exec_query_stats or other related views.

20. Fixing Orphaned Users:

Orphaned user User1 in the SURESHDB database.

When we run **sp_change_users_login** with the **REPORT** option, we can see ~~that it is an orphaned~~that an orphaned user.

```
EXEC sp_change_users_login 'REPORT'
```

UserName	UserSID

User1	0xA5B5548F3DC81D4693E769631629CE1D

To fix this orphaned user all we have to do is run **sp_change_users_login** with the **UPDATE_ONE** action and tell SQL Server the name ~~of the orphaned~~of orphaned user and the name of the appropriate login.

```
EXEC sp_change_users_login 'UPDATE_ONE','User1','User1'
```

